Basic Overview of Machine Learning Methods

JEFF GILL

Distinguished Professor

Department of Government, Department of Mathematics & Statistics

Member, Center for Neuroscience and Behavior

Founding Director, Center for Data Science

American University

Background: What is Machine Learning?

▶ Machine Learning is a collection of tools mostly for classification and prediction.

▶ Most of these you already know or are close to something you already know, and the vocabulary is simply different (logit).

▶ The focus is mostly on prediction, regression, and classification.

▶ The term is not as new as one would think (Samuel, *IBM Journal of Research and Development*, 1959).

▶ Modern definition: "A computer program is said to learn from experience E with regard to some class of tasks T and performance measure P, if its performance at tasks T, as measured by P, improves with experience E." (Mitchell, *Machine Learning*, 1997).

▶ Common applications: credit card data analysis, speech recognition, text analysis, fraud detection, self-driving cars, website ads, and many more.

▶ A lot of these applications were previously addressed with rigid rule-based systems.

Supervised Learning

▶ There is a target variable Y that we want to predict given feature variables $\mathbf{X}$ by learning a function such that $F(\mathbf{X}) = Y$.

▶ When Y is interval measured this is regression and when Y is categorical this is classification.

▶ A key goal is to find the *best* F possible to estimate/predict future (unseen) data.

▶ This is a different than classical statistical inference that focuses on the data at hand.

Unsupervised Learning

▶ Here there is an identified target Y.

▶ The goal is to identify groupings within the data in different ways.

▶ Cluster identification, principle components analysis, are the classic examples.

▶ Note that this dichotomy is not strict and there are lots of tools in between supervised and unsupervised learning: weakly supervised and hybrids/combinations.

More Differences from Classic Statistical Inference

▶ In most social science settings the first emphasis is on selecting a parsimonious set of control variables ("under the horizon") and set of theory-based variables ("over the horizon").

▶ But in ML the typical strategy is to have start with many explaining variables and reduce the number with a hold-out/test strategy to winnow the number down possibly.

▶ There usually is not a concern about how many are left, unlike regular statistical inference (AIC, BIC, DIC, etc.).

▶ A primary reason for this is that ML is most often used with big data so the $p \gg n$ is unlikely to be a concern.

▶ A lot of the work in ML is done to "process" the data with transformations of the data as part of the fitting process and then have a relatively simple $F()$.

▶ This is different trend than statistics.

Vocabulary (this is important!)

▶ Features: explanatory variables.

▶ Labels: outcome variables.

▶ Examples: the subjects (data).

▶ Learning: training with the data, estimating a function, building a model (the process).

▶ Deep Learning: a hierarchical process wherein complex representations (models) are created in the algorithm from simple representations in a dynamic process ($F()$ is created from combining many far simpler $f()$ functions in multilevel structure).

▶ Overfitting: a sin in ML, the model is too closely aligned with a single dataset, including it's error component, ruining generality (application to future datasets).

▶ Underfitting: the model does not explain the underlying phenomenon.

Clustering

▶ The most common unsupervised learning method: asserting/inferring substantive groups in the data cases.

▶ This is hardly new/modern, e.g. statisticians have been arguing about the number of clusters in the "galaxy" data for at least 50 years.

▶ Most often the determination of clusters is done spatially with respect to fixed data points based on a distance measure like Euclidean, Manhattan, or Mahalanobis.

▶ Not always so though, e.g. text analysis, genetics, social network analysis.

▶ Two different settings: users specifies the number of clusters in advance (relatively easy), or the number of clusters is not known in advance (relatively hard).

▶ We often get different cluster arrangements with different algorithms.

Example of Cluster Configurations

Partitions for $n = 4$

$p = 1$	$p = 2$	$p = 3$	$p = 4$
$\mathfrak{R}_1$	$\mathfrak{R}_2$	$\mathfrak{R}_3$	$\mathfrak{R}_4$

| $y_1 y_2 y_3 y_4$ | $y_1 \mid y_2 y_3 y_4$ $y_2 \mid y_1 y_3 y_4$ $y_3 \mid y_1 y_2 y_4$ $y_4 \mid y_1 y_2 y_3$ | $y_1 y_2 \mid y_3 y_4$ $y_1 y_3 \mid y_2 y_4$ $y_1 y_4 \mid y_2 y_3$ | $y_1 \mid y_2 \mid y_3 y_4$ $y_1 \mid y_3 \mid y_2 y_4$ $y_1 \mid y_4 \mid y_2 y_3$ $y_2 \mid y_3 \mid y_1 y_4$ $y_2 \mid y_4 \mid y_1 y_3$ $y_3 \mid y_4 \mid y_1 y_2$ | $y_1 \mid y_2 \mid y_3 \mid y_4$ |

▶ Four Cluster Classes, Five Configuration Classes

▶ The number of configuration classes in each cluster class is $b(n, p)$

▷ $b(n, p)$ = partitions of the integer n into p components ≥ 1

▷ $b(4, 1) = 1$, $b(4, 2) = 2$, $b(4, 3) = 1$, $b(4, 4) = 1$,

Notes On Cluster Configurations

▶ So for this $n = 4$ illustration in each of the cluster classes, $p = (1, 2, 3, 4)$, there are: $b(n, p) = (1, 2, 1, 1)$ configuration classes, and $(1, 7, 6, 1)$ partition types.

▶ The number of partition types, for a given n and p is a Stirling number of the second kind from:

$$\left\{ \begin{matrix} n \\ p \end{matrix} \right\} = \frac{1}{p!} \sum_{j=0}^{p} (-1)^{p-j} \binom{p}{j} j^n.$$

▶ In the example there are 15 total possible partitions (models), the Bell number for $n = 4$ from:

$$B_n = \frac{1}{e} \sum_{j=0}^{\infty} \frac{j^n}{j!}$$

▶ We connect these because a Bell number can be expressed as the sum of Stirling numbers of the second kind:

$$B_n = \sum_{p=0}^{n} \left\{ \begin{matrix} n \\ p \end{matrix} \right\}$$

▶ For a fixed m, the number of configuration classes $b(n, m)$ grows as $\frac{n^{m-1}}{m!(m-1)!}$ with increasing n.

A Typology of Clustering

▶ Agglomerative Clustering: at the beginning each data point is its own cluster then the algorithm combines the points into clusters (e.g. K-Means).

▶ Divisive: at the beginning all the data points are in the same cluster then the algorithm breaks the apart into a number of clusters (e.g. Mean Shift).

▶ We can actually get very different cluster configurations based on which algorithm is used.

K-Means Clustering Algorithm

▶ This the oldest and most common method.

▶ It is basically a variance minimizer in an ANOVA sense (sum of squares).

▶ Guaranteed to converge (always "works").

▶ You have to have the number of clusters in advance.

▶ Process:

    ▷ distribute k centroids in the data space, randomly, uniformly, or purposefully
    ▷ each data point gets assigned to the nearest centroid creating clusters
    ▷ within each cluster move the centroid to the spatial mean of the data points
    ▷ repeat until the centroids do not move anymore.

▶ Users often do this process multiple times with different starting points to gain confidence.

▶ K-Means is very fast and therefore very useful with big data.

K-Means in R

```r
# GET FISHER'S IRIS DATA
data(iris)
head(iris, 10)
   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1           5.1         3.5          1.4         0.2  setosa
2           4.9         3.0          1.4         0.2  setosa
3           4.7         3.2          1.3         0.2  setosa
4           4.6         3.1          1.5         0.2  setosa
5           5.0         3.6          1.4         0.2  setosa
6           5.4         3.9          1.7         0.4  setosa
7           4.6         3.4          1.4         0.3  setosa
8           5.0         3.4          1.5         0.2  setosa
9           4.4         2.9          1.4         0.2  setosa
10          4.9         3.1          1.5         0.1  setosa

# REMOVE SPECIES LABEL
iris2 <- iris[,-5]
```

K-Means in R

```
# LOAD PACKAGES
library(ClusterR); library(cluster)

# RUN K-MEANS, centers is k, nstarts is the number of random starts
set.seed(1234)
( k.means.out <- kmeans(x=iris2, centers = 3, nstart = 25) )

K-means clustering with 3 clusters of sizes 50, 62, 38

Cluster means:
  Sepal.Length Sepal.Width Petal.Length Petal.Width
1     5.006000    3.428000     1.462000    0.246000
2     5.901613    2.748387     4.393548    1.433871
3     6.850000    3.073684     5.742105    2.071053
```

K-Means in R

```
Clustering vector:
  [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
 [35] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 3 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
 [69] 2 2 2 2 2 2 2 2 2 3 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 3 2
[103] 3 3 3 3 2 3 3 3 3 3 3 2 2 3 3 3 3 2 3 2 3 2 3 3 2 2 3 3 3 3 3 2 3 3
[137] 3 3 2 3 3 3 2 3 3 3 2 3 3 2

Within cluster sum of squares by cluster:
[1] 15.15100 39.82097 23.87947
 (between_SS / total_SS =  88.4 %)

# CONFUSION MATRIX
table(iris$Species, k.means.out$cluster)

             1  2  3
  setosa    50  0  0
  versicolor 0 48  2
  virginica  0 14 36
```

K-Means in R

```
# PLOT 1
postscript("CLASSES/Class.Big.Data/k.means.plot1.ps")
par(mfrow=c(1,1),mar=c(5,5,2,2),lwd=2,col.axis="white",col.lab="white",
    col.main="white", col.sub="white", col="white",bg="slategray", cex.lab=1.3)
plot(iris2[c("Sepal.Length", "Sepal.Width")], col = k.means.out$cluster,
     main = "K-Means Assignments", pch="+")
points(k.means.out$centers[, c("Sepal.Length", "Sepal.Width")],
       col = 1:3, pch = 19, cex = 2)
dev.off()
```

K-Means Plot 1

K-Means in R

```
# PLOT 2
postscript("CLASSES/Class.Big.Data/k.means.plot2.ps")
clusplot(iris2[, c("Sepal.Length", "Sepal.Width")], k.means.out$cluster,
        lines = 1, shade = TRUE, color = TRUE, labels = 2, plotchar = TRUE,
        span = FALSE, main = paste("K-Means Assignments"),
        xlab = "Sepal.Length", ylab = "Sepal.Width")
dev.off()
```

K-Means Plot 2
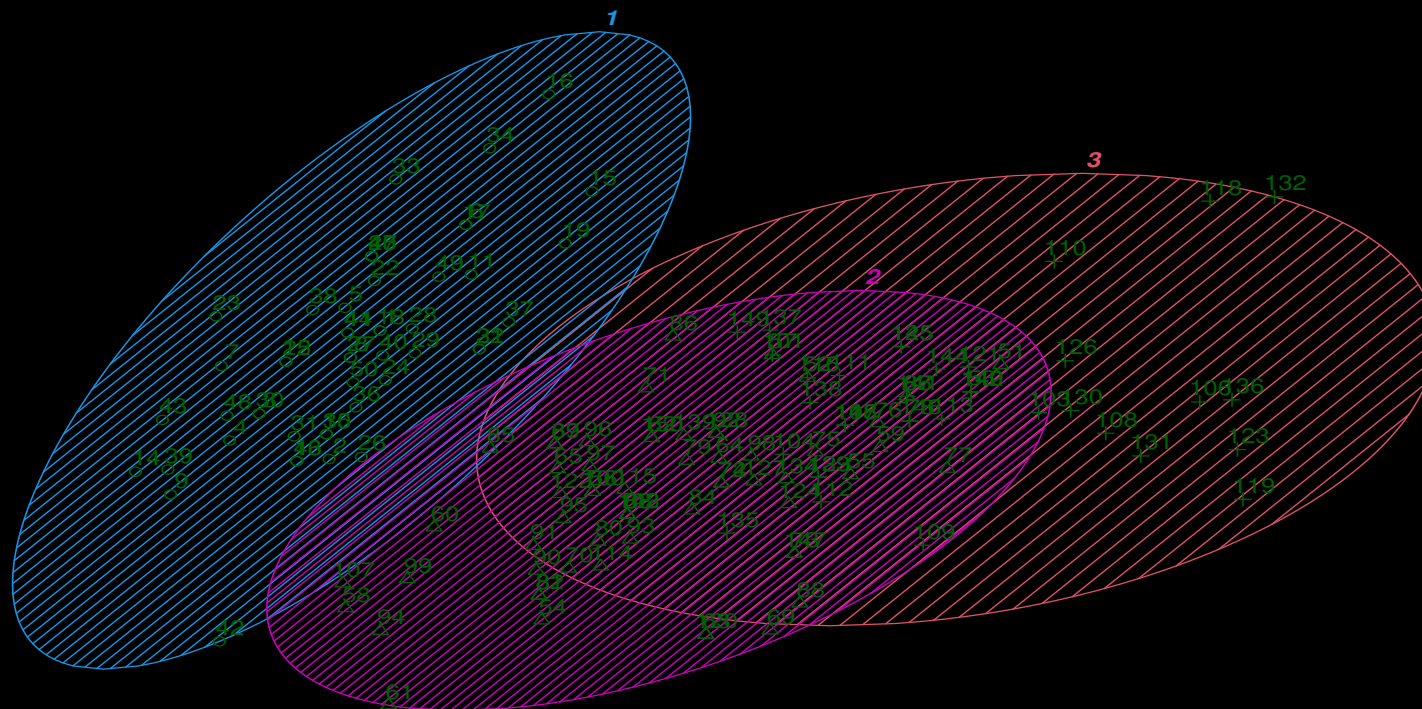
Limitations of K-Means

▶ The need to choose k in advance.

▶ A heavy dependence on initial centroid locations.

▶ Performs poorly when real clusters are of very different sizes as it makes cluster sizes equal.

▶ Sensitivity to spatial outliers.

▶ Does poorly in high dimensions.

▶ Assumes that the variance within each cluster is the same.

▶ Will find clusters even if there aren't any in the data.

Mean Shift Clustering

▶ A somewhat different approach that starts with the points rather than the centroids and is based on density.

▶ Instead of k we have to define a bandwidth parameter h.

▶ Algorithm:

▷ create a circular window around every data point of radius h so every data point defines a cluster
▷ get the mean position of the points inside each window, make this a centroid
▷ move the center of the window to this centroid
▷ repeat until convergence.

▶ Eventually all points in the same cluster will end up with basically the same steps towards convergence.

▶ This is conditional on a reasonable bandwidth value, which must be chosen carefully.

▶ So when multiple windows overlap the algorithm keeps only the window containing the most points.

Mean Shift in R

```
library(meanShiftR)
( ms.out <- meanShift(queryData = as.matrix(iris2),trainData = as.matrix(iris2),
    algorithm="LINEAR") )
t(ms.out$assignment)
```

```
     [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14] [,15] [,16] [,17] [,18] [,19] [,20]
[1,]    1    1    1    1    1    1    1    1    1     1     1     1     1     1     1     1     1     1     1     1
     [,21] [,22] [,23] [,24] [,25] [,26] [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38] [,39]
[1,]     1     1     1     1     1     1     1     1     1     1     1     1     1     1     1     1     1     1     1
     [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50] [,51] [,52] [,53] [,54] [,55] [,56] [,57] [,58]
[1,]     1     1     1     1     1     1     1     1     1     1     1     2     2     2     2     2     2     2     2
     [,59] [,60] [,61] [,62] [,63] [,64] [,65] [,66] [,67] [,68] [,69] [,70] [,71] [,72] [,73] [,74] [,75] [,76] [,77]
[1,]     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2
     [,78] [,79] [,80] [,81] [,82] [,83] [,84] [,85] [,86] [,87] [,88] [,89] [,90] [,91] [,92] [,93] [,94] [,95] [,96]
[1,]     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2     2
     [,97] [,98] [,99] [,100] [,101] [,102] [,103] [,104] [,105] [,106] [,107] [,108] [,109] [,110] [,111] [,112]
[1,]     2     2     2     2      2      2      2      2      2      2      2      2      2      2      2      2
     [,113] [,114] [,115] [,116] [,117] [,118] [,119] [,120] [,121] [,122] [,123] [,124] [,125] [,126] [,127] [,128]
[1,]      2      2      2      2      2      2      2      2      2      2      2      2      2      2      2      2
     [,129] [,130] [,131] [,132] [,133] [,134] [,135] [,136] [,137] [,138] [,139] [,140] [,141] [,142] [,143] [,144]
[1,]      2      2      2      2      2      2      2      2      2      2      2      2      2      2      2      2
     [,145] [,146] [,147] [,148] [,149] [,150]
[1,]      2      2      2      2      2      2
```

Mean Shift in R

```
# PLOT
postscript("CLASSES/Class.Big.Data/mean.shift.plot.ps")
par(mfrow=c(1,1),mar=c(5,5,2,2),lwd=2,col.axis="white",col.lab="white",
    col.main="white", col.sub="white", col="white",bg="slategray", cex.lab=1.3)
plot(iris2[c("Sepal.Length", "Sepal.Width")][ms.out$assignment==1,], col = "white",
     xlim=c(4,8), ylim=c(2,4.5),main = "Mean Shift Assignments", pch="+")
points(iris2[c("Sepal.Length", "Sepal.Width")][ms.out$assignment==2,], col = "red",
    pch="+")
points(ms.out$value[c(1,150),], col = c("white","red"), pch = 19, cex = 2)
dev.off()
```

Mean Shift Plot

Limitations of Mean Shift

▶ Due to the cycling through the points and then cycling through the clusters on each iteration it can be very compute-intensive, especially with big data.

▶ The choice of h is critical: too small and convergence may not happen, too large and distinct actual clusters in the data get be merged.

▶ Cluster distinctions are "sharp" and not overlapping, which makes this method popular in computer vision, but that may be less appropriate in the social sciences.

Spectral Clustering

▶ SC works very differently than what we've seen so far by valuing "connectivity" over convex boundaries.

▶ SC can find non-convex clusters.

▶ It uses a graph (network) setup where the data are nodes and the edges are similarity.

▶ Algorithm:

    ▷ Create the similarity matrix from the distance between each pair of data points: which is square with values $n(n-1)/2$ that we care about.

    ▷ Perform *graph partitioning* whereby the edges between clusters are given low weights and the edges with clusters are given high weights.

    ▷ Compute the eigenvalues and eigenvectors.

    ▷ The k eigenvectors are treated as data and supplied to k-means to create the clusters for the original data.

▶ SC performs well but is slower with big data.

▶ There are also some important selection parameters and algorithmic decisions to be made.

Spectral Clustering Technical Details

► We have n points $x_.$ in $\mathfrak{R}^k$, where p is the dimension of the data (two so far in the slides).

► And d_{ij} is the distance between x_i and x_j (generally Euclidean but any type).

► As a similarity graph $\mathbf{S}$ use the *radial-kernel grain matrix*, whose elements are defined by:

$$s_{ij} = \exp\left[\frac{-d_{ij}}{c}\right]$$

where positive c is a (selection) scale parameter.

► Alternatively we could use as a similarity graph the *mutual K-nearest neighbor graph* that starts with defining N_K as the symmetric set of nearby points: s_{ij} is positive in N_K if i is in the *K-nearest neighbors* of j. Then assemble the sets of nearest neighbors and assign them the edge weight w_{ij} and the excluded relationships are assigned zero.

► There are many other alternatives.

Spectral Clustering Technical Details

▶ Generally the result of the similarity process with its edge weights is called an adjacency matrix and denoted $\mathbf{W}$ with elements w_{ij}.

▶ Vertex i has *degree* that is the sum of its weights:

$$g_i = \sum_{i=1:n, \neg i} w_{ij}.$$

▶ Now define the diagonal matrix G that collects $g_i, i = 1 \dots n$.

▶ The *standardized graph Laplacian* is given by:

$$\mathbf{L} = \mathbf{I} - \mathbf{G}^{-1}\mathbf{W}$$

▶ Now find the m smallest eigenvectors (selection parameter) of $\mathbf{L}$ corresponding to the m smallest eigenvalues: $\mathbf{E}_{n \times m}$.

▶ Cluster the rows of $\mathbf{E}_{n \times m}$ with K-Means (or some alternative).

Eigen-Analysis of Matrices

▶ Every $p \times p$ matrix $\mathbf{X}$ has p scalar values, $\lambda_i, i = 1, \ldots, p$, such that

$$\mathbf{X}\mathbf{e}_i = \lambda_i \mathbf{e}_i$$

for some corresponding vector $\mathbf{e}_i$.

▶ In this *decomposition*, λ_i is called an eigenvalue of $\mathbf{X}$ and $\mathbf{e}_i$ is called an *eigenvector* of $\mathbf{X}$.

▶ These eigenvectors are linearly independent.

▶ These are also called the *characteristic roots* and *characteristic vectors* of $\mathbf{X}$, and the process is also called *spectral decomposition*.

▶ The full *eigendecomposition* of the original square matrix is given by: $\mathbf{X} = \mathbf{E}(\lambda \mathbf{I})\mathbf{E}^{-1}$ where $\mathbf{E}$ is a matrix with the eigenvectors down columns.

▶ The eigenvalues and eigenvectors are only guaranteed to be real-valued if the original square matrix is symmetric.

▶ The *characteristic equation* is given by: $\text{diag}((\mathbf{X} - \boldsymbol{\lambda}\mathbf{I})\mathbf{E}) = \mathbf{0}$.

▶ The eigenvalues and eigenvectors are found by solving the characteristic equation.

Basic Eigenanalysis

▶ A symmetric square matrix $\mathbf{X}$ is given by

$$\mathbf{X} = \begin{bmatrix} 1.000 & 0.880 & 0.619 \\ 0.880 & 1.000 & 0.716 \\ 0.619 & 0.716 & 1.000 \end{bmatrix}.$$

▶ Using R (note the descending order):

```
eigen( matrix(c(1,0.88,0.619,0.88,1,0.716,0.619,0.716,1),3,3) )
eigen() decomposition
$values
[1] 2.4820708 0.4100160 0.1079132

$vectors
           [,1]        [,2]       [,3]
[1,] 0.5850593  0.5127477  0.6283274
[2,] 0.6071390  0.2367286 -0.7585128
[3,] 0.5376688 -0.8252571  0.1728089
```

▶ Note: in the eigenvector matrix returned by R the eigenvectors are the *columns*.

Eigenanalysis Uniqueness

▶ Eigenvalues and eigenvectors are associated.

▶ For each eigenvector of a given matrix $\mathbf{X}$ there is exactly one corresponding eigenvalue such that

$$\lambda = \frac{\mathbf{e}'\mathbf{X}\mathbf{e}}{\mathbf{e}'\mathbf{e}}.$$

▶ For each eigenvalue of the matrix there is an infinite number of eigenvectors, all determined by scalar multiplication.

▶ If $\mathbf{h}$ is an eigenvector corresponding to the eigenvalue λ, then $s\mathbf{h}$ is also an eigenvector corresponding to this same eigenvalue where s is any nonzero scalar.

Eigenanalysis General Properties

▶ The number of nonzero eigenvalues is the rank of the $\mathbf{X}$.

▶ The sum of the eigenvalues is the trace of $\mathbf{X}$.

▶ The product of the eigenvalues is the determinant of $\mathbf{X}$.

▶ A matrix is singular if and only if it has a zero eigenvalue, and the rank of the matrix is the number of nonzero eigenvalues.

▶ If there are no zero-value eigenvalues, then the eigenvectors determine a basis for the space determined by the size of the matrix ($\mathfrak{R}^2$, $\mathfrak{R}^3$, etc.).

▶ Symmetric nonsingular matrices have eigenvectors that are perpendicular (orthogonal).

Spectral Clustering in R

```r
library("speccalt")
kern.partitioning <- local.rbfdot(iris2) # USES A KDE TO DEFINE SIMILARITY THRESHOLD
speccalt(kern.partitioning) # AUTOMATIC CHOICE OF m CLUSTERS
```

```
  [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2
 [61] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
[121] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```r
speccalt(kern.partitioning,5) # CHOOSE w
```

```
  [1] 2 5 5 5 2 2 2 2 5 5 2 2 5 5 2 2 2 2 2 2 2 2 2 2 2 5 2 2 2 5 5 2 2 2 5 5 2 2 5 2 2 1 5 2 2 5 2 5 2 2 3 3 3 1 3 1 3 1 3 1
 [61] 1 3 1 3 1 3 1 1 1 1 3 3 3 3 3 3 3 3 3 1 1 1 1 3 1 3 3 1 1 1 1 3 1 1 1 1 1 3 1 1 4 3 4 3 4 4 1 4 3 4 4 3 4 3 3 4 3 4 4 1
[121] 4 3 4 3 4 4 3 3 3 4 4 4 3 3 3 4 4 3 3 4 4 4 3 4 4 4 3 3 4 3
```

Contrived Difficult Case

Contrived Difficult Case

```
ruler <- seq(1.56,1.58,length=500)
y1 <- tan(ruler)/1000 + rnorm(500,0,1.0)
y1[y1 > 20] <- 10; y1[y1 < -20] <- -10
synth.dat <- cbind(ruler,y1)
kern.partitioning <- local.rbfdot(synth.dat)
sc.fit <- speccalt(kern.partitioning,2)

postscript("CLASSES/Class.Big.Data/spectral.sim.ps")
par(mfrow=c(1,1),mar=c(5,5,2,2),lwd=2,col.axis="white",col.lab="white",
    col.main="white", col.sub="white", col="white",bg="slategray", cex.lab=1.3)
plot(synth.dat[1:270,],pch="+",col="black", xlim=c(1.560,1.580), ylim=c(-15,12),
    xlab="Explanatory Variable",ylab="Outcome Variable")
points(synth.dat[271:500,],pch="+",col="red")
dev.off()
```

Principal Components Analysis

▶ Principal components analysis (PCA) is a means of data reduction with big data through rotation in the sample space of observations.

▶ This is very useful in big data analysis where $p \gg n$.

▶ PCA re-expresses the variability of the data such that the total amount of variance is preserved but:

▷ Axes are enumerated in descending order of variance explained. That is, the first dimension explains the most variance, the second dimension explains the second-most variance, and so on.

▷ The new axes are uncorrelated with each other: they are orthogonal.

▷ If there exists correlation in the original data then it is expressed as zero length along some dimensions after the rotation of axes.

Principal Components Analysis

- ▶ If there are p explanatory variables in the original data and there is some correlation between these variables, then PCA produces a rotation in the p dimensions except that some number of these, q, will be of zero length, where the magnitude of q indicates the extent of the correlation between variables and $p - q$ indicates the extent of orthogonal information in the data.

- ▶ Therefore if all p variables are uncorrelated then the axes are already orthogonal and there is no need to perform PCA.

- ▶ Conversely, if the p variables are perfectly correlated then there exists only one dimension worth of information in the data and all but one of the axes will have data of zero-length after PCA.

Principal Components Analysis

▶ Consider for a moment only two variables: X_1 and X_2, which are assumed for simplicity to have mean zero each and unit variance.

▶ With this variance assumption the correlation reduces to covariance.

▶ If the correlation between these two variables is actually zero, then the equiprobability contours (concentric lines indicating equal probability of occurrence) of these two variables is circular.

▶ On the other hand, if there is a non-zero ρ value then the shape of the equiprobability contours will be elliptical where the cosine of the angle of intersection from the longest elliptical axis to the original x-axis (measured at the origin since zero mean is assumed for both variables) is equal to ρ.

▶ In the extreme case of perfect correlation between X_1 and X_2 the equiprobability contours condense to a single line.

Illustration of 2-D PCA

From https://setosa.io/ev/principal-component-analysis/.

Principal Components Analysis

▶ Begin with an $n \times p$ data matrix $\mathbf{X}$, where variables are organized in columns, and standardize.

▶ Define $\mathbf{R}$ as the correlation matrix corresponding to $\mathbf{X}$ along with a matrix $\mathbf{E}$ of the eigenvectors of the $\mathbf{R}$ matrix with the constraint that squared rows and columns of $\mathbf{E}$ sum to one.

▶ Then by standard spectral theory (Lax 1997, Chapter 6), the matrix defined by:

$$\boldsymbol{\lambda} = \mathbf{E}'\mathbf{R}\mathbf{E}$$

is a matrix containing the descending eigenvalues along the diagonal and zeros elsewhere.

▶ It is in fact the variance-covariance matrix of the rotation defined by the principal components.

▶ So each eigenvalue, $\lambda_1, \lambda_2, \ldots$, is the variance of a principal component where the first principal component now accounts for the largest variance by construction.

▶ Furthermore, since the off-diagonal elements are all zero, then the correlation has been removed in the new coordinate system.

Principal Components Analysis

▶ The importance of the $\mathbf{E}$ matrix is that it provides the transformation of the data points from the original metric to the PCA metric through simple matrix multiplication: $\mathbf{Y} = \mathbf{XE}$.

▶ Thus $\mathbf{Y}$ are the points in the new rotated coordinate system where the variance structure is preserved: the principal component scores.

▶ The usefulness of this transformation is that it is one-to-one and therefore reversible, $\mathbf{X} = \mathbf{E}'\mathbf{Y}$ because of the orthogonal property of the $\mathbf{E}$ matrix.

▶ The $\mathbf{E}$ matrix of normalized eigenvectors is orthogonal, meaning that:

$$\mathbf{E}'\mathbf{E} = \mathbf{E}\mathbf{E}' = \mathbf{I}.$$

▶ Also this property also allows us to modify $\boldsymbol{\lambda} = \mathbf{E}'\mathbf{R}\mathbf{E}$ according to:

$$\boldsymbol{\lambda}\mathbf{E}' = \mathbf{E}'\mathbf{R}\mathbf{E}\mathbf{E}' = \mathbf{E}'\mathbf{R}$$

Principal Components Analysis

▶ We can also define a new matrix called the component loadings according to:

$$\mathbf{L} = \mathbf{E}\boldsymbol{\lambda}^{\frac{1}{2}}$$

where the square root on $\boldsymbol{\lambda}$ is simply the square root of each diagonal element.

▶ The $\mathbf{L}$ matrix is theoretically important due to two related multiplicative properties:

$$\begin{aligned} \mathbf{LL}' &= \mathbf{E}\boldsymbol{\lambda}^{\frac{1}{2}}(\mathbf{E}\boldsymbol{\lambda}^{\frac{1}{2}})' \\ &= \mathbf{E}\boldsymbol{\lambda}\mathbf{E}' \\ &= \mathbf{E}(\mathbf{E}'\mathbf{R}\mathbf{E})\mathbf{E}' \\ &= \mathbf{R} \end{aligned} \qquad \begin{aligned} \mathbf{L}'\mathbf{L} &= (\mathbf{E}\boldsymbol{\lambda}^{\frac{1}{2}})'\mathbf{E}\boldsymbol{\lambda}^{\frac{1}{2}} \\ &= (\boldsymbol{\lambda}^{\frac{1}{2}})'\mathbf{E}'\mathbf{E}\boldsymbol{\lambda}^{\frac{1}{2}} \\ &= \boldsymbol{\lambda} \end{aligned}$$

▶ So the product of component loadings is either equal to the correlation matrix ($\mathbf{R}$) or the diagonal eigenvalue matrix ($\boldsymbol{\lambda}$), depending on the order of matrix multiplication.

Principal Components Analysis, Worked Example

▶ Consider the following contrived dataset and its standardized column variables:

$$\tilde{\mathbf{X}} = \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 3 \\ 3 & 1 & 2 \\ 4 & 3 & 6 \\ 5 & 5 & 5 \\ 6 & 7 & 6 \\ 7 & 9 & 9 \\ 8 & 8 & 8 \\ 9 & 8 & 3 \end{bmatrix} \qquad \mathbf{X} = \begin{bmatrix} -1.461 & -1.109 & -1.385 \\ -1.095 & -0.421 & -0.652 \\ -0.730 & -1.453 & -1.018 \\ -0.365 & -0.765 & 0.448 \\ 0.000 & -0.076 & 0.081 \\ 0.365 & 0.612 & 0.448 \\ 0.730 & 1.300 & 1.547 \\ 1.095 & 0.956 & 1.181 \\ 1.461 & 0.956 & -0.652 \end{bmatrix}$$

```
tilde.X <- matrix(c(1,2,1,2,4,3,3,1,2,4,3,6,5,5,5,6,7,6,7,9,9,8,8,8,9,8,3),
    ncol=3,byrow=TRUE)
X <- scale(X)
```

Principal Components Analysis, Worked Example

▶ The correlation matrix from $\mathbf{X}$ is:

$$\mathbf{R} = \begin{bmatrix} 1.000 & 0.880 & 0.619 \\ 0.880 & 1.000 & 0.716 \\ 0.619 & 0.716 & 1.000 \end{bmatrix}.$$

The eigenvalues and eigenvectors are found by solving the characteristic equation: $|\mathbf{R} - \boldsymbol{\lambda}| = 0$.

▶ This produces the matrices:

$$\mathbf{E} = \begin{bmatrix} -0.585 & -0.514 & 0.628 \\ -0.607 & -0.236 & -0.759 \\ -0.538 & 0.825 & 0.174 \end{bmatrix}, \qquad \boldsymbol{\lambda} = \begin{bmatrix} 2.482 & 0.00 & 0.000 \\ 0.000 & 0.41 & 0.000 \\ 0.000 & 0.00 & 0.108 \end{bmatrix}.$$

```
R <- cor(X)
E <- eigen(R)$vectors
lambda <- eigen(R)$values * diag(3)
L <- E %*% chol(lambda)
```

Principal Components Analysis, Worked Example

▶ This means that the proportion of the total variance explained by each of the three principal components are $2.482/3 = 0.827, 0.41/3 = 0.137, 0.108/3 = 0.036$.

▶ By the same reasoning, the second principal component explains 13.7% of the total variance and the third principal component explains 3.6% of the total variance.

▶ The component scores are produced by pre-multiplying the original data matrix by $\mathbf{E}$, producing:

$$\mathbf{Y} = \begin{bmatrix} 2.272 & -0.130 & -0.316 \\ 1.247 & 0.124 & -0.482 \\ 1.857 & -0.122 & 0.467 \\ 0.437 & 0.737 & 0.429 \\ 0.003 & 0.085 & 0.072 \\ -0.826 & 0.038 & -0.157 \\ -2.049 & 0.595 & -0.259 \\ -1.856 & 0.186 & 0.168 \\ -1.084 & -1.513 & 0.078 \end{bmatrix}.$$

```
diag(lambda)/3
Y <- X %*% E
```

Principal Components Analysis, Worked Example

▶ Here the mean for each $\mathbf{Y}$ variable (the columns) remains zero and the corresponding variance is no longer unity but rather the corresponding diagonal value of $\boldsymbol{\lambda}$.

▶ Because of these component scores result from the simple matrix multiplication defined in $\mathbf{Y} = \mathbf{XE}$, they are in fact linear combinations of the original data with weights determined by the eigenvector matrix $\mathbf{E}$.

▶ For instance, the first value of $\mathbf{Y}$ is produced from:

$$Y_{11} = \sum_{j=1}^{3} E_{.1} X_{1.}$$
$$= 0.585 \times -1.461 + 0.607 \times -1.109 + 0.538 \times -1.385$$
$$= 2.272.$$

```
Y11 <- sum(E[,1] %*% X[1,])
```

Association Rules

▶ This is a data mining tool (unsupervised) that answers the question of finding whether certain "items" occur together more frequently than randomness predicts.

▶ The general setup is $X_1 \ldots X_k \longrightarrow Y$ having support S and confident C.

▶ So when a purchase/event contains $X_1 \ldots X_k$ at least $C\%$ of the time then $\mathbf{Y}$ also occurs, and there are at least $S\%$ of these transactions out of those observed.

▶ So we care about frequency and statistical significance of these associations.

▶ For two individual items X and Y *lift* is the ratio of the probability of joint occurrence over the product of individual occurrences under the assumption of independence:

$$\text{lift} = \frac{p(X,Y)}{p(X)p(Y)}.$$

▶ So a value near one implies that the items are not associated, small values less than one imply that these rarely occur together, and large values imply common joint occurrence.

▶ There are then two elements: Antecedent (if) this is an item/group of items that are found in the itemsets, and Consequent (then): these are later associated with an Antecedent or set of

Antecedents.

Association Rules

▶ Commonly association rules are applied to big commercial datasets: $p \approx 10^5, n \approx 10^10$.

▶ Algorithm:

▷ Tabulate all combinations of items in a dataset that occur together with a minimum frequency: frequent itemsets.

▷ Assert association rules that parameterize co-occurrence within the frequent itemsets.

▶ This is implemented with rule mining algorithms that identify a basket of items $X_1 \ldots X_k$ relative to Y.

▶ Most often the Xs are binary such as bought/didn't buy, and the dataset is checkout summaries so that x_{ij} is person i's realized (observed) purchase status of item j.

Association Rules

▶ Often the events do not happen often enough for a single unit of study so *regions* of the event space are considered: groups of events, events over time, groups of people, etc.

▶ Define S_j as the support of events for the jth variable and $s_j \subseteq S_j$ as a subset.

▶ We want to find a subset of variable values such that the intersection has relatively large probability:

$$\text{conjunctive rule} = p \left[\bigcap_{j=1}^{p} (X_j \in s_j) \right].$$

▶ Only 2 types of subsets are considered due to data size frequently encountered:

    ▷ s_j is a single value of X_j denoted $s_j = \eta_{j0}$
    ▷ all values that X_j can assume: $s_j = S_j$.

NEED FIGURE HERE

Association Example: Market Basket Analysis

▶ Consider a dataset of 9,409 questionnaires filled out by shoppers in the SF Bay Area where the demographic questions only are studied:

Feature Number	Feature	Categories	Type
1	Sex	2	Categorical
2	Marital Status	5	Categorical
3	Age	7	Ordinal
4	Education	6	Ordinal
5	Occupation	9	Categorical
6	Income	9	Ordinal
7	Years in Bay Area	5	Ordinal
8	Dual Income	3	Categorical
9	Number in Household	9	Ordinal
10	Number of Children	9	Ordinal
11	Householder Status	3	Categorical
12	Type of Home	5	Categorical
13	Ethnic Classification	8	Categorical
14	Language in Home	3	Categorical

Association Example: Market Basket Analysis

▶ There are missing data, which was case-wise deleted by Hastie, Tibshirani, and Friedman.

▶ They used software called `Apriori` by Christian Borgelt.

▶ Ordinal features were cut at their median to produce dichotomous features (dummy variables).

▶ Categorical variables given a treatment contrast.

▶ The final dataset after pre-processing was $6{,}875 \times 50$.

▶ `Apriori` found 6,288 association rules that had 5 or less predictors and support of at least 10%.

Association Example: Market Basket Analysis

▶ **Association Rule # 1:** Support 25%, Confidence 97%, Lift 1.03

Antecedents	Consequent
Number in Household = 1	
Number of Children = 0	Language in Home = English

▶ **Association Rule # 2:** Support 13.4%, Confidence 80.8%, Lift 2.13

Antecedents	Consequent
Language in Home = English	
Householder Status = own	
Occupation = Professional/Managerial	Income $\geq$ \$40,000

▶ **Association Rule # 3:** Support 26.5%, Confidence 82.8%, Lift 2.15

Antecedents	Consequent
Language in Home = English	
Income ¡ \$40,000	
Marital Status = Not Married	
Number of Children = 0	Education <u>not</u> college graduate or graduate study

K-Nearest Neighbors

▶ This is a classifier in the family called *memory-based models*.

▶ It is very simple and very fast spatial approach, even with very large or very fast data.

▶ For the basic method there are only two Decisions: k and the distance metric (Euclidean, Manhattan, etc.).

▶ With k users usually do some trial and error since too small means high variance and too large may miss important local features.

▶ Basic algorithm:

  ▷ Identify a multidimensional query point x_0.
  ▷ Find the k points nearest to this points: x_r for $r = 1, \ldots, k$, usually using Euclidean distance in feature space: $d(r) = \|x_{(r)} - x_0\|$. ($\|\boldsymbol{v}\| = (v_1^2 + v_2^2 + \cdots + v_n^2)^{\frac{1}{2}} = (\boldsymbol{v}' \cdot \boldsymbol{v})^{\frac{1}{2}}$).
  ▷ For the feature of interest assign an attribute to x_o based on some "voting" criteria, usually majority rule (ties settled at random).
  ▷ Repeat as new points are identified.

K-Nearest Neighbors

▶ Typically the variables (dimensions) are standardized to mean zero and standard deviation one such that the measurement of variables does not lead to domination/sublimation.

▶ This is a old and popular tool that is well-suited to big data problems, including EKG patterns, handwriting analysis, image analysis, satellite data, internet traffic, political ideology, and more.

▶ It performs well even when the decision barrier is very irregular.

▶ There are some challenges when the features are both continuous and categorical since more decisions need to be made.

▶ Other challenges include sparsity and very high dimensions.

▶ There are many, many extensions/enhancements like weighting.

K-Nearest Neighbors, Example

```
library(neighbr)
data(iris); head(iris)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species ID
1          5.1         3.5          1.4         0.2  setosa  1
2          4.9         3.0          1.4         0.2  setosa  2
3          4.7         3.2          1.3         0.2  setosa  3
4          4.6         3.1          1.5         0.2  setosa  4
5          5.0         3.6          1.4         0.2  setosa  5
6          5.4         3.9          1.7         0.4  setosa  6

iris$ID <- c(1:150) # APPEND AN ID NUMBER COLUMN

train_set <- iris[1:145,] # USE FIRST 145 CASES AS TRAINING WITH ALL FEATURES

test_set <- iris[146:150,-c(4,5,6)] # REMOVE PREDICTED VARIABLES FROM TEST DATA
```

K-Nearest Neighbors, Example

```
fit <- knn(train_set=train_set,test_set=test_set, k=5, categorical_target="Species",
    continuous_target= "Petal.Width", comparison_measure="euclidean",
    return_ranked_neighbors=3, id="ID")
fit$test_set_scores
```

	categorical_target	continuous_target	neighbor1	neighbor2	neighbor3
146	virginica	2.04	78	142	140
147	virginica	1.64	73	124	134
148	virginica	1.86	111	116	78
149	virginica	2.06	137	116	138
150	virginica	1.90	115	128	84

K-Nearest Neighbors, Example

```
fit <- knn(train_set=train_set,test_set=test_set, k=9, categorical_target="Species",
    continuous_target= "Petal.Width", comparison_measure="euclidean",
    return_ranked_neighbors=9, id="ID")
fit$test_set_scores
```

	categorical_target	continuous_target	neighbor1	neighbor2	neighbor3	neighbor4
146	virginica	2.022222	78	142	140	111
147	virginica	1.577778	73	124	134	84
148	virginica	1.933333	111	116	78	117
149	virginica	2.133333	137	116	138	111
150	virginica	1.855556	115	128	84	139

	neighbor5	neighbor6	neighbor7	neighbor8	neighbor9
146	113	117	116	53	141
147	127	112	120	74	64
148	134	112	138	113	142
149	117	104	125	145	101
150	102	143	71	122	134

Support Vector Machines

▶ This is a very useful and very popular supervised classifier on interval measured data.

▶ Consider multidimensional, interval-measured data where we want to separate (classify) points by specifying lines in each dimension which collectively determine a hyperplane.

▶ The idea is to get the best separation possible and to not have the line too close to data points to increase generalizability with future data.

▶ This is done by finding the hyperplane maximizes the *margin* of the training data.

▶ Unlike many other classifiers SVM does not provide probabilities.

SVM Illustration

Support Vector Machines

▶ Training sample Points are called *support vectors* and those closest to the margin are the most influential.

▶ For non-linear structures in the features kernels (e.g. radial basis function) provide an efficient tool for separation (not discussed here).

▶ The training data consists of n pairs: $\{(x_1, y_1) \ldots (x_n, y_n)\}$, where $x_i \in \mathfrak{R}^p$ and $y_i \in \{-1, 1\}$.

▶ Define a hyperplane by:
$$x : f(x) = \beta_0 + x\beta = 0$$
where β is of length 1.

▶ If the points are separable then $y_i f(x_i) > 0 \ \forall i$.

▶ This means that we can find the hyperplane that provides the maximum margin between the $y = -1$ group and the $y = 1$ group:
$$\max_{\beta_0, \beta} = M \quad \text{subject to} \quad y_i(f(x) = \beta_0 + x\beta) \geq M, \quad \forall i.$$

▶ So the band is M away from the hyperplane in both directions.

Support Vector Machines

▶ If there is overlap in the feature space that cause some points to be on the wrong side of any specified hyperplane then specify slack variables, $\xi_i \geq 0$, $\sum \xi_i \leq C$, for some specified constant C, and we now specify:

$$\max_{\beta_0, \beta} = M(1 - \mathbf{x}_i) \quad \text{subject to} \quad y_i(f(x) = \beta_0 + x\beta) \geq M, \quad \forall i.$$

▶ Since miscalculations occur when $\mathbf{x}_i > 0$ then C bounds the total number of training misclassifications.

▶ With overlap we drop the norm constraint on β and now define $M = 1/||\beta||$.

▶ This leads to a modified procedure:

$$\min ||\beta|| \quad \text{subject to} \begin{cases} y_i(f(x) = \beta_0 + x\beta) \geq 1 = \xi_i, \quad \forall i \\ \xi_i \geq \sum \mathbf{x}_i \leq C \end{cases}$$

▶

SVM Example in R

```
set.seed(999)
x <- matrix(rnorm(60), 30, 2)
y <- rep(c(-1, 1), c(15,15))
x[y==1,] = x[y==1,] + 1
par(mfrow=c(1,1),mar=c(5,5,2,2),lwd=2,col.axis="white",col.lab="white",
    col.main="white", col.sub="white", col="white",bg="slategray", cex.lab=1.3)
plot(x, col = y+51, pch = "+", cex=2)

library(e1071)
svm.example <- data.frame(x, y = factor(y))
svm.out <- svm(y ~ ., data = svm.example, kernel = "linear", cost = 5, scale = FALSE)
```

SVM Example in R

```
summary(svm.out)


Parameters:
   SVM-Type:  C-classification
 SVM-Kernel:  linear
       cost:  5
Number of Support Vectors:  9
 ( 4 5 )
Number of Classes:  2
Levels:
 -1 1


plot(svm.out,svm.example)
```

SVM Default Graph

SVM Example in R

```
Beta.times.Label <- svm.out$coefs
cbind(svm.out$SV,Beta.times.Label)  # THE 9 SUPPORT VECTORS
          X1         X2
1  -0.2817402 2.3826642  5.000000
11  2.3254637 0.5143633  5.000000
12  1.1339774 1.0084981  5.000000
15  1.9576504 1.3006654  5.000000
17  3.0683351 0.9491223 -1.601791
21  0.7714367 1.8833359 -5.000000
25  0.8747315 1.9331901 -3.398209
28  1.1151596 0.7207841 -5.000000
30  1.8733210 1.4344990 -5.000000
```

SVM Example in R

```
svm.pred <- fitted(svm.out) # TEST WITH TRAINING DATA
rbind(svm.pred[1:15],svm.pred[16:30])
      1 2 3 4 5 6 7 8 9 10 11 12 13 14 15
[1,]  2 1 1 1 1 1 1 1 1  1  1  1  1  1  2
[2,]  2 2 2 2 2 2 2 2 2  2  2  2  1  2  2
Levels: -1 1


table(svm.pred,y) # ACCURACY SUMMARY
         y
svm.pred -1   1
      -1 13   1
       1  2  14
```

Bootstrapping for Standard Errors

▶ **Big Idea**: sometimes it is difficult to get the sampling properties of an estimator, even a commonly used one.

▶ Some statistics have known variance properties for finite samples and some do not. Does this mean we should only use the former unless we have population data?

▶ Definitive citations: Efron (1979), Efron and Tibshirani (1993).

▶ **Case Study:** suppose we have a dataset on leukemia, ignoring a whole host of things and condensing our analysis down to two variables: *CD4 Count/10* a dichotomous outcome indicating that there was a relapse from a remission stage:

Relapse	94	197	16	38	99	141	23		
No Relapse	52	104	146	10	50	31	40	27	46

▶ Note that these data are *imbalanced*.

Bootstrapping for Standard Errors

▶ The question is whether there is a difference by CD4 count, and the natural choice of test is the difference of means: $\bar{x}_{\text{relapse}} = 86.86, \ \bar{x}_{\text{no relapse}} = 56.22$.

▶ This is easy since we know that:

$$SE(\bar{x}_{\text{relapse}}) = \sqrt{(s^2_{\text{relapse}}/n_{\text{relapse}})} = 25.24,$$
$$SE(\bar{x}_{\text{no relapse}}) = \sqrt{(s^2_{\text{no relapse}}/n_{\text{no relapse}})} = 14.14$$

▶ But we also know that the mean is not very resistant to outliers and it could be that a notable case, and one could be driving the subsequent findings.

▶ So what about using the median instead of the mean? This is obvious choice in one sense, but it leaves us with no closed form solution for the standard error.

Bootstrapping for Standard Errors

▶ So consider the following algorithm, for some statistic of interest, θ:

1. Draw B "bootstrap" samples of size n, independently, **with replacement** from the sample $\mathbf{x}$:

$$\mathbf{x}^{*1}, \mathbf{x}^{*2}, \ldots, \mathbf{x}^{*B}$$

(note the notation to differentiate the bootstrap sample from the original sample).

2. Calculate the sample statistic of interest, θ^{*b} for each bootstrap sample, and the mean of these statistics:

$$\bar{\theta}^* = \frac{1}{B} \sum_{b=1}^{B} \theta^{*b}$$

3. Estimate the bootstrap standard error of the statistic by:

$$\mathsf{Var}(\theta) = \frac{1}{B-1} \sum_{b=1}^{B} \left(\theta^{*b} - \bar{\theta}^* \right)^2$$

where obviously $SE(\theta) = \sqrt{\mathsf{Var}(\theta)}$.

▶ We call the limit of this standard error as B goes to infinity is called the *ideal* bootstrap estimate, and this procedure is called the *nonparametric* bootstrap estimate.

Bootstrapping for Standard Errors

```r
relapse <- c(94,197,16,38,99,141,23)
no.relapse <- c(52,104,146,10,50,31,40,27,46)
B <- 1000
no.relapse.mat <- relapse.mat <- NULL
for (i in 1:B)  {
    relapse.mat    <- rbind(relapse.mat,   sample(relapse,length(relapse),
                           replace=TRUE))
    no.relapse.mat <- rbind(no.relapse.mat,sample(no.relapse,length(no.relapse),
                           replace=TRUE))
}

relapse.mean <- mean(apply(relapse.mat,1,mean))
relapse.se <- sqrt(var(apply(relapse.mat,1,mean)))
no.relapse.mean <- mean(apply(no.relapse.mat,1,mean))
no.relapse.se <- sqrt(var(apply(no.relapse.mat,1,mean)))
relapse.median <- mean(apply(relapse.mat,1,median))
relapse.median.se <- sqrt(var(apply(relapse.mat,1,median)))
no.relapse.median <- mean(apply(no.relapse.mat,1,median))
no.relapse.median.se <- sqrt(var(apply(no.relapse.mat,1,median)))
```

Bootstrapping for Standard Errors

```
final.relapse.mat <- rbind( c(relapse.mean, relapse.se,
                              no.relapse.mean, no.relapse.se),
    c(relapse.median, relapse.median.se, no.relapse.median, no.relapse.median.se) )
dimnames(final.relapse.mat) <-
    list( c("Mean","Median"), c("Relapse Est","Relapse SE",
                               "No Relapse Est","No Relapse SE") )


final.relapse.mat
       Relapse Est  Relapse SE  No Relapse Est  No Relapse SE
Mean    88.19143    21.57597        54.86444       12.45360
Median  84.37000    36.40226        44.07000       12.65810
```

Bootstrapping for Standard Errors

```
par(mfrow=c(2,2),bg="white")
hist(apply(relapse.mat,1,mean), freq=FALSE, main="Bootstrap, Relapse Mean")
ruler <- seq(20,160,length=100)
lines(ruler,dnorm(ruler,relapse.mean,relapse.se),lwd=3)
hist(apply(relapse.mat,1,median), main="Bootstrap, Relapse Median")
hist(apply(no.relapse.mat,1,mean), ylim=c(0,0.04),freq=FALSE,
    main="Bootstrap, No Relapse Mean")
ruler <- seq(20,100,length=100)
lines(ruler,dnorm(ruler,no.relapse.mean,no.relapse.se),lwd=3)
hist(apply(no.relapse.mat,1,median), main="Bootstrap, No Relapse Median")
```

Bootstrapping for Standard Errors

Random Forests

▶ Suppose we have a dataset with N cases and M features and we want to classify.

▶ General Algorithm:

  ▷ Draw B bootstrap samples of size n, independently, **with replacement** from the data. The size of n depends on the context of the problem.

  ▷ For each of these B train a decision tree sampling $m \ll M$ features uniformly from the full M and the collection of decision trees picks the best features given the features that they individually have.

  ▷ New data is tested with all of the decision trees and the final result is an aggregation such as majority vote.

▶ RFs can handle large problems very easily.

▶ Notice that this process can be parallelized for computational efficiency.

▶ RFs also provide a proximity matrix showing similarity between all of the points by counting the proportion of times two selected data points are classified together at the bottom of the trees.

Bagging: Bootstrap Aggregation

▶ This is more general *ensemble* method than random forests since any ML process can be used.

▶ Draw B bootstrap samples of size n, independently, **with replacement** from the data. The size of n depends on the context of the problem and the size of N.

▶ Train a model on each of these B datasets.

▶ Obtain a out-of-sample test dataset in the typical fashion and use each of the B models to predict.

▶ The procedure can also easily be parallelized.

Boosting

▶ This is another very general procedure whereby one iteratively trains classifiers using the data cases where the previous model produced misclassifications.

▶ Therefore each iteration gets a smaller dataset than all of the previous models.

▶ Algorithm:

  ▷ Give all the data points equal weights.
  ▷ For the jth model in the series:
    ▷ train a classifier using the current weights
    ▷ predict from the training data
    ▷ determine the error from this prediction
    ▷ calculate new weights based on the errors in the jth classification
  ▷ Repeat.
  ▷ At the end of the run a weighted average of the predictions from all of the models where the weight is proportional to each of the model's prediction accuracy.

Final Thoughts on Machine Learning

▶ There are usually a lot of decisions to be made that contextual to the data and the problem:

  ▷ size of the training set versus the test set within the same data set or outside of it.

  ▷ tuning parameters

  ▷ validation measures such as categorical prediction comparisons, distance measures, and cross-validation.

  ▷ leakage: the modeling process does not have access to the training anymore at testing time.